

Working with (big) data

Big Data for Public Policy

Christoph Goessmann


Law, Economics, and Data Science Group (Prof. Elliott Ash)

Course: 860-0033-00L Big Data for Public Policy, Spring 2023

2 March 2023

Survey

ETH Edu App



Web app

<https://eduapp-app1.ethz.ch/>



iOS



Android

Projects

Choice

- Research Paper
 - Sergio Galletta (sergio.galletta@gess.ethz.ch)
- Web App
 - Christoph Goessmann (christoph.goessmann@gess.ethz.ch)

Tentative Outline

Technical Part


- How to efficiently work with computers?
 - Computer architecture (I/O, CPU, GPU)
 - Controlling a computer (command line, ssh)
 - Shapes and forms of data (text vs. binary, compression)
 - Reproducibility and automation (virtual environments, logging, git, etc.)
 - Data acquisition (scrapping)
- APIs and web apps
- Machine Learning

Questions?

- Interrupt with questions / comments anytime!
- Happy to adapt the lectures to your needs.

Computer Architecture

Typical data flow



CPU bound (CPU load ~ 100%)

- increase number of CPUs/cores (parallelization)
- use GPUs instead of CPUs
- change algorithm (e.g., profiling)


I/O bound (CPU load << 100%)

- reduce I/O operations
- switch to faster memory


Computer Architecture

I/O bound vs CPU bound processes


I/O bound



Somewhere in between



CPU bound




Computer Architecture

Common data types

Type	Size	Range
boolean	1 byte	0,1
smallint	2 bytes	-32,768 to +32,767
int	4 bytes	-2,147,483,648 to +2,147,483,647
bigint	8 bytes	-9,223,372,036,854,775,808 to +9,223,372,036,854,775,807
real / single precision float	4 bytes	6 decimal digits precision
double precision float	8 bytes	15 decimal digits precision

Computer Architecture

Operating System



Linux/Unix dominant

- 100% of all TOP500 supercomputers [1]
- > 80% of public servers (web, mail, DNS, etc.) on the internet [2]
- Some packages (e.g., CuPy, PyTorch ROCm) are only fully supported on Linux

[1] <https://top500.org/statistics/overtime/>, accessed 28 February 2023

[2] <https://w3techs.com/>, accessed 28 February 2023

Command Line

Basic commands

\$ **pwd** (print working directory)

\$ **cd** *path* (change directory)

\$ **mkdir** *newd* (make directory)

\$ **ll** / **ls -al** (list contents of working directory)

\$ **cat** *file* (display content of file)

\$ **head -n file** (display first n lines of file)

\$ **tail -n file** (display last n lines of file)

\$ **rm** *file* (remove/delete file)

\$ **com > file** (write output of command com to file)

\$ **man** *command* (display command's manual)

\$ **CTRL + C** (abort command)

\$ **CTRL + D** (end of stream)

Command Line

Text editor

Create new file / edit existing one

```
$ vim file.txt
```

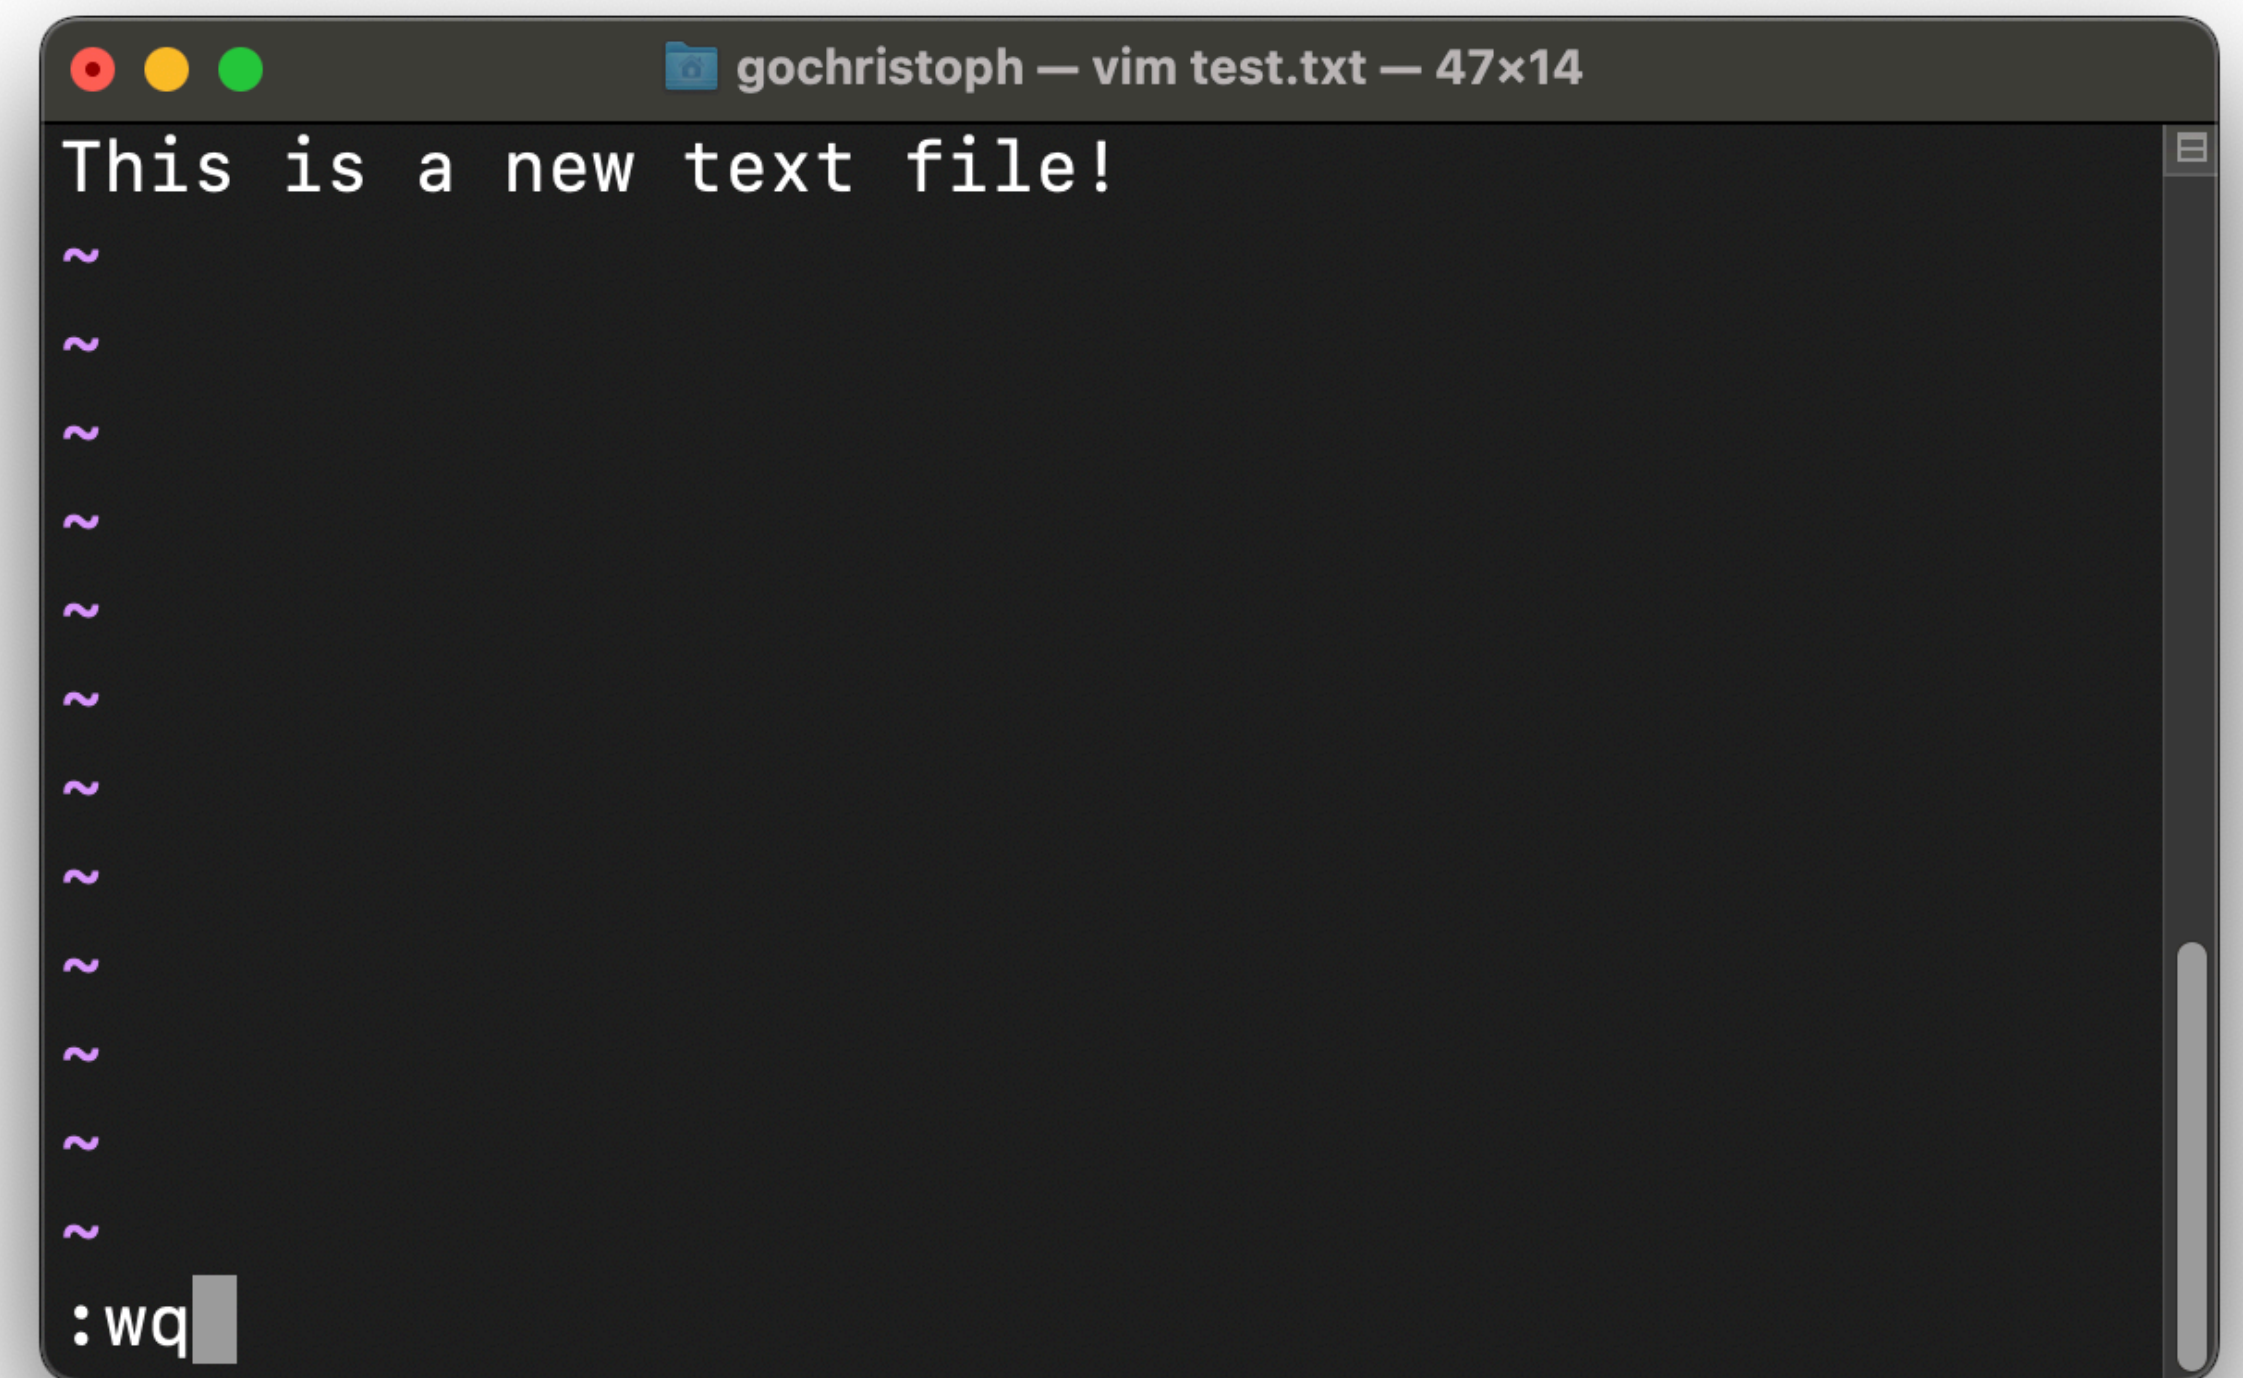
Command mode
(default at start, invoke with ESC):

Switch to insert mode: **i**

Save: **:w** *+Enter*

Save and quit: **:wq** *+Enter*

Quit without saving: **:q!** *+Enter*



Command Line

Monitoring CPU load & profiling

Task manager

```
$ htop / top
```

Average CPU load of your script

```
$ time python my_program.py (local)
```

```
$ myjobs -j jobid (ETH Euler cluster)
```

Profiling with cProfile and SnakeViz

```
$ python -m cProfile -o my_program.prof my_program.py
```

```
$ snakeviz my_program.prof
```



Command Line

Virtual environments with pipenv

Install pipenv

```
$ pip install pipenv
```

Change to project directory and install a package

```
$ cd myproject
```

```
$ pipenv install package
```

List of packages and dependencies is written to Pipfile/Pipfile.lock.

Whenever you are in the project directory, you have two options to **use the virtual environment**:

```
$ pipenv shell
```

```
$ pipenv run python  
myscript.py
```

Use pyenv for managing different python versions on the same machine.

Hands-on

Example

10 million documents

- Each document i has a 128-dim. feature vector $a_{i,*}$ (e.g., LDA topics).
- Want to calculate all cosine similarities between documents i and j :

$$s_{i,j} = \frac{a_{i,*} \cdot a_{j,*}}{\|a_{i,*}\| \|a_{j,*}\|}$$

(naively corresponds to matrix multiplication of $1e7 \times 128$ matrix with its transpose)

Feature matrix ($1e7 \times 128$):

$$(a_{i,j}) = \begin{bmatrix} a_{1,1} & a_{1,2} & a_{1,3} & \dots & a_{1,128} \\ a_{2,1} & a_{2,2} & a_{2,3} & \dots & a_{2,128} \\ a_{3,1} & a_{3,2} & a_{3,3} & \dots & a_{3,128} \\ \dots & \dots & \dots & \dots & \dots \\ a_{10^7,1} & a_{10^7,2} & a_{10^7,3} & \dots & a_{10^7,128} \end{bmatrix}$$

Similarity matrix ($1e7 \times 1e7$):

$$(s_{i,j}) = (a_{i,j})(a_{i,j})^T$$

Hands-on

First considerations

Document vector (1 x 128):

$$a_{i,*} = [a_{i,1} \quad a_{i,2} \quad a_{i,3} \quad \dots \quad a_{i,128}]$$

→ size = $1e7 * 8 \text{ byte} = \mathbf{80 \text{ MB}}$

Feature matrix (1e7 x 128):

$$(a_{i,j}) = \begin{bmatrix} a_{1,1} & a_{1,2} & a_{1,3} & \dots & a_{1,128} \\ a_{2,1} & a_{2,2} & a_{2,3} & \dots & a_{2,128} \\ a_{3,1} & a_{3,2} & a_{3,3} & \dots & a_{3,128} \\ \dots & \dots & \dots & \dots & \dots \\ a_{10^7,1} & a_{10^7,2} & a_{10^7,3} & \dots & a_{10^7,128} \end{bmatrix}$$

→ size = $1e7 * 128 * 8 \text{ byte} = \mathbf{10.2 \text{ GB}}$

Similarity matrix (1e7 x 1e7):

$$(s_{i,j}) = (a_{i,j})(a_{i,j})^T$$

→ size = $1e7 * 128 * 8 \text{ byte} = \mathbf{800 \text{ TB}}$

Is it realistic to calculate all cosine similarities?

Hands-on

1. Create a virtual environment

```
$ pip install pipenv  
$ mkdir bdpp_session1  
$ cd bdpp_session1  
$ pipenv install numpy  
$ pipenv install snakeviz  
$ pipenv shell
```

2. Write a python script (**matrix_math_test.py**) that:

- creates a random feature matrix
- computes the cosine similarity between the feature matrix and 100 randomly drawn rows

Hands-on

3. Time your script:

```
$ time python my_program.py
```


Extra: export OPENBLAS_NUM_THREADS=**1/0**

4. Profile your script with cProfile and SnakeViz

```
$ python -m cProfile -o my_program.prof my_program.py
```

```
$ snakeviz my_program.prof
```

Hands-on



Hands-on

Extras on Euler (Slurm & GPUs)

CPU only

```
#!/bin/bash
```

```
#SBATCH -n 1
```

```
#SBATCH --cpus-per-task=1
```

```
#SBATCH --time=4:00:00
```

```
#SBATCH --mem-per-cpu=24576
```

```
#SBATCH --mail-type=ALL
```

```
python matrix_math_test.py
```

```
$ sbatch ...
```

```
$ myjobs / seff jobid
```

Runtime: 120s

GPU (numpy -> cupy):

```
#!/bin/bash
```

```
#SBATCH --gpus=1
```

```
#SBATCH --gres=gpumem:22G
```

```
#SBATCH --time=4:00:00
```

```
#SBATCH --mail-type=ALL
```


```
module load gcc/8.2.0 python_gpu && python  
matrix_math_test_gpu.py
```

```
$ sbatch ...
```

```
$ myjobs / seff jobid
```


Runtime: 5s (24 x faster)

Data



End-of-Lecture Survey

ETH Edu App



Web app

<https://eduapp-app1.ethz.ch/>



iOS



Android